

> Initial setup

↳ 3 cells hidden

> Data pre-processing

↳ 8 cells hidden

> A/B Testing with Baseline Model

↳ 18 cells hidden

Best current model is that with augmentation and
using the full image

```
def set_seed(seed: int = 42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False
seed=42
set_seed(42)

SAVE_DIR = "/content/bike_lane_runs"
os.makedirs(SAVE_DIR, exist_ok=True)
print("Saving checkpoints/configs to:", SAVE_DIR)
```

Saving checkpoints/configs to: /content/bike_lane_runs

```
# updated old function
```

```
def train_one_model(
    name,
    use_lane_crop,
    max_epochs=15,
    lr=1e-4,
    batch_size=8,
    save_dir=SAVE_DIR,
    seed=seed,
    min_delta=1e-3,           # min improvement in val_acc to count as be-
    patience=3                 # stop if no meaningful improvement after 3
):
    print(f"Training model: {name} (use_lane_crop={use_lane_crop})")

    set_seed(seed)

    train_loader, val_loader, test_loader = make_loaders(
        use_lane_crop, batch_size=batch_size
    )

    backbone = make_backbone()
    model = SegFormerBlockedClassifier(backbone, num_classes=2).to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.AdamW(model.parameters(), lr=lr)

    train_losses = []
    val_accs = []

    best_val_acc = -1.0
    best_epoch = -1
    epochs_no_improve = 0

    os.makedirs(save_dir, exist_ok=True)
    best_ckpt_path = os.path.join(save_dir, f"{name}_best.pt")

    for epoch in range(1, max_epochs + 1):
        model.train()
        running_loss = 0.0

        for pixel_values, labels in train_loader:
            pixel_values = pixel_values.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()
            logits = model(pixel_values)
```

```
loss = criterion(logits, labels)
loss.backward()
optimizer.step()

running_loss += loss.item() * labels.size(0)

train_loss = running_loss / len(train_loader.dataset)
val_acc, _, _ = evaluate(model, val_loader)

train_losses.append(train_loss)
val_accs.append(val_acc)

print(f"[{name}] Epoch {epoch}/{max_epochs} - "
      f"train_loss: {train_loss:.4f} - val_acc: {val_acc:.4f}")

# check for improvement
if val_acc > best_val_acc + min_delta:
    best_val_acc = val_acc
    best_epoch = epoch
    epochs_no_improve = 0

    torch.save(
        {
            "model_state_dict": model.state_dict(),
            "optimizer_state_dict": optimizer.state_dict(),
            "epoch": best_epoch,
            "val_acc": best_val_acc,
            "seed": seed,
            "use_lane_crop": use_lane_crop,
            "lr": lr,
            "batch_size": batch_size,
            "min_delta": min_delta,
            "patience": patience,
        },
        best_ckpt_path,
    )
    print(f" ↴ New best val_acc={best_val_acc:.4f} at epoch {best_epoch} - "
          f"saved to {best_ckpt_path}")
else:
    epochs_no_improve += 1
    print(f" ↴ No significant improvement (Δval_acc={val_acc - "
          f"epochs_no_improve={epochs_no_improve}/{patience}})")

# early stopping?
if epochs_no_improve >= patience:
```

```
print(f"\nEarly stopping triggered at epoch {epoch} "
      f"(no improvement > {min_delta} for {patience} epochs"
      break

# plot
epochs_ran = range(1, len(train_losses) + 1)

plt.figure(figsize=(8, 4))
plt.plot(epochs_ran, train_losses, marker="o")
plt.title(f"{name} - Training Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.grid(True)
plt.show()

plt.figure(figsize=(8, 4))
plt.plot(epochs_ran, val_accs, marker="o")
plt.title(f"{name} - Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.ylim(0, 1.0)
plt.grid(True)
plt.show()

# best checkpoint
print(f"\n[{name}] Loading best checkpoint from epoch {best_epoch}:"
      f"(val_acc={best_val_acc:.4f})")
checkpoint = torch.load(best_ckpt_path, map_location=device)
model.load_state_dict(checkpoint["model_state_dict"])

# test eval
test_acc, test_cm, test_report = evaluate(model, test_loader)
print(f"\n[{name}] TEST accuracy: {test_acc:.4f}")
print(f"[{name}] TEST confusion matrix (rows = true, cols = pred")
print(test_cm)
print(f"\n[{name}] TEST classification report:")
print(test_report)

config = {
    "name": name,
    "seed": seed,
    "max_epochs": max_epochs,
    "lr": lr,
    "batch_size": batch_size,
    "use_lane_crop": use_lane_crop,
```

```

        "augment": True,
        "backbone": "nvidia/segformer-b0-finetuned-ade-512-512",
        "best_epoch": best_epoch,
        "best_val_acc": float(best_val_acc),
        "test_accuracy": float(test_acc),
        "min_delta": float(min_delta),
        "patience": patience,
    }

    config_path = os.path.join(save_dir, f"{name}_config.json")
    with open(config_path, "w") as f:
        json.dump(config, f, indent=2)
    print(f"\n[{name}] Saved run config to {config_path}")

    history = {
        "train_losses": train_losses,
        "val_accs": val_accs,
    }

    return model, (test_acc, test_cm, test_report, history)

```

In this cell, I train the baseline segformers classifier on full-frame images w data augmentation and early stopping to get a clean model run:

- min_delta = 0.001 (val accuracy must improve by at least 0.001)
- patience = 3 (stop if no such improvement for 3 consecutive epochs)

```

baseline_model, baseline_metrics = train_one_model(
    name="baseline_full_frame_aug_es",
    use_lane_crop=False,           # using full images
    max_epochs=15,
    lr=1e-4,
    batch_size=8,
    save_dir=SAVE_DIR,
    seed=seed,
    min_delta=1e-3,               # require at least 0.001 val_acc
    patience=3                    # stop after 3 epochs with no su
)

```



```


Training model: baseline_full_frame_aug_es (use_lane_crop=False)
/content/split_camera/train: found 2739 images (use_lane_crop=False, au
/content/split_camera/val: found 1175 images (use_lane_crop=False, augm
/content/split_camera/test: found 38 images (use_lane_crop=False, augme

```

```
[baseline_full_frame_aug_es] Epoch 1/15 - train_loss: 0.5365 - val_acc: 0.8502
↳ New best val_acc=0.8502 at epoch 1, saved to /content/bike_lane_rur
[baseline_full_frame_aug_es] Epoch 2/15 - train_loss: 0.3386 - val_acc: 0.8732
↳ New best val_acc=0.8732 at epoch 2, saved to /content/bike_lane_rur
[baseline_full_frame_aug_es] Epoch 3/15 - train_loss: 0.2789 - val_acc: 0.9149
↳ New best val_acc=0.9149 at epoch 3, saved to /content/bike_lane_rur
[baseline_full_frame_aug_es] Epoch 4/15 - train_loss: 0.2657 - val_acc: 0.9174
↳ New best val_acc=0.9174 at epoch 4, saved to /content/bike_lane_rur
[baseline_full_frame_aug_es] Epoch 5/15 - train_loss: 0.2288 - val_acc: 0.9174
↳ No significant improvement ( $\Delta$ val_acc=-0.0068), epochs_no_improve=1
[baseline_full_frame_aug_es] Epoch 6/15 - train_loss: 0.2322 - val_acc: 0.9174
↳ No significant improvement ( $\Delta$ val_acc=-0.0255), epochs_no_improve=2
[baseline_full_frame_aug_es] Epoch 7/15 - train_loss: 0.2344 - val_acc: 0.9174
↳ No significant improvement ( $\Delta$ val_acc=-0.0068), epochs_no_improve=3
```

Early stopping triggered at epoch 7 (no improvement > 0.001 for 3 epochs)





```
[baseline_full_frame_aug_es] Loading best checkpoint from epoch 4 (val_
[baseline_full_frame_aug_es] TEST accuracy: 0.3947
[baseline_full_frame_aug_es] TEST confusion matrix (rows = true, cols =
[[ 4 23]
 [ 0 11]]

[baseline_full_frame_aug_es] TEST classification report:
      precision    recall   f1-score   support
notblocked       1.000     0.148     0.258      27
      blocked       0.324     1.000     0.489      11

      accuracy          0.395      38
macro avg       0.662     0.574     0.373      38
weighted avg     0.804     0.395     0.325      38



[baseline_full_frame_aug_es] Saved run config to /content/bike_lane_rur
```


```
baseline_modelC, baseline_metricsC = train_one_model(
    name="baseline_lanecropped_frame_aug_es",
    use_lane_crop=True,                      # using full images
    max_epochs=15,
    lr=1e-4,
    batch_size=8,
    save_dir=SAVE_DIR,
    seed=seed,
    min_delta=1e-3,                         # require at least 0.001 val_ac
    patience=3                               # stop after 3 epochs with no !
)
```

```
Training model: baseline_lanecropped_frame_aug_es (use_lane_crop=True)
/content/split camera/train: found 2739 images (use lane crop=True, aug
```

```
/content/split_camera/val: found 1175 images (use_lane_crop=True, augme
/content/split_camera/test: found 38 images (use_lane_crop=True, augmen
[baseline_lanecropped_frame_aug_es] Epoch 1/15 - train_loss: 0.5284 - v
    ↴ New best val_acc=0.8460 at epoch 1, saved to /content/bike_lane_run
[baseline_lanecropped_frame_aug_es] Epoch 2/15 - train_loss: 0.3397 - v
    ↴ New best val_acc=0.8664 at epoch 2, saved to /content/bike_lane_run
[baseline_lanecropped_frame_aug_es] Epoch 3/15 - train_loss: 0.2890 - v
    ↴ New best val_acc=0.8902 at epoch 3, saved to /content/bike_lane_run
[baseline_lanecropped_frame_aug_es] Epoch 4/15 - train_loss: 0.2536 - v
    ↴ No significant improvement ( $\Delta$ val_acc=-0.0026), epochs_no_improve=1/
[baseline_lanecropped_frame_aug_es] Epoch 5/15 - train_loss: 0.2453 - v
    ↴ New best val_acc=0.9115 at epoch 5, saved to /content/bike_lane_run
[baseline_lanecropped_frame_aug_es] Epoch 6/15 - train_loss: 0.2304 - v
    ↴ New best val_acc=0.9149 at epoch 6, saved to /content/bike_lane_run
[baseline_lanecropped_frame_aug_es] Epoch 7/15 - train_loss: 0.2085 - v
    ↴ No significant improvement ( $\Delta$ val_acc=-0.0077), epochs_no_improve=1/
[baseline_lanecropped_frame_aug_es] Epoch 8/15 - train_loss: 0.2165 - v
    ↴ No significant improvement ( $\Delta$ val_acc=-0.0094), epochs_no_improve=2/
[baseline_lanecropped_frame_aug_es] Epoch 9/15 - train_loss: 0.2016 - v
    ↴ No significant improvement ( $\Delta$ val_acc=0.0009), epochs_no_improve=3/3
```

Early stopping triggered at epoch 9 (no improvement > 0.001 for 3 epochs)





```
[baseline_lanecropped_frame_aug_es] Loading best checkpoint from epoch  
[baseline_lanecropped_frame_aug_es] TEST accuracy: 0.2895  
[baseline_lanecropped_frame_aug_es] TEST confusion matrix (rows = true,  
[[ 0 27]  
 [ 0 11]]  
  
[baseline_lanecropped_frame_aug_es] TEST classification report:  
          precision    recall   f1-score   support  
  
notblocked      0.000     0.000     0.000      27  
  blocked       0.289     1.000     0.449      11  
  
accuracy           -         -         -      38  
macro avg       0.145     0.500     0.224      38  
weighted avg     0.084     0.289     0.130      38
```

```
[baseline_lanecropped_frame_aug_es] Saved run config to /content/bike_1  
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result)  
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result)  
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result)
```

During training, we see early stopping is triggered at epoch 7, meaning that validation accuracy stopped improving meaningfully after epoch 4. Epoch 4's validation accuracy = 0.9174

When I evaluate it on the held-out cam68 test set,

- Test accuracy: 0.3947
- For notblocked (27 images):
 - 4 correctly predicted as notblocked
 - 23 incorrectly predicted as blocked
 - → very low recall (0.148), but precision = 1.0 (when the model says "notblocked", it's always correct, just very rare)
- For blocked (11 images):
 - 11 correctly predicted as blocked
 - 0 missed
 - → recall = 1.0, but precision is only 0.324 (many notblocked images are being flagged as blocked)

Overall, this operating point is extremely conservative in the sense of rarely missing a blocked lane (100% blocked recall), but it generates many false alarms on notblocked images, leading to low overall accuracy (~ 39.5%). This motivates the threshold-tuning step that follows.

▼ Tuned classifier

```
# use_lane_crop=False bc cropped was worse
train_loader, val_loader, test_loader = make_loaders(
    use_lane_crop=False,
    batch_size=8
)
/content/split_camera/train: found 2739 images (use_lane_crop=False, a
/content/split_camera/val: found 1175 images (use_lane_crop=False, augm
/content/split_camera/test: found 38 images (use_lane_crop=False, augme
```

```
# collect p(blocked) and true labels (probabilities)
def get_probs_and_labels(model, loader):
    model.eval()
    all_probs = []
    all_labels = []

    with torch.no_grad():
        for pixel_values, labels in loader:
            pixel_values = pixel_values.to(device)
            labels = labels.to(device)

            logits = model(pixel_values) # (B, 2)
            probs = torch.softmax(logits, dim=1) # (B, 2)
            p_blocked = probs[:, 1] # prob of class 1

            all_probs.extend(p_blocked.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    return np.array(all_probs), np.array(all_labels)

val_probs, val_labels = get_probs_and_labels(baseline_model, val_loader)
test_probs, test_labels = get_probs_and_labels(baseline_model, test_loader)

print("Val shape:", val_probs.shape, val_labels.shape)
print("Test shape:", test_probs.shape, test_labels.shape)

# val_probs[i] = model's probability that image i is blocked
# val_labels[i] = 0 (notblocked) or 1 (blocked)
```

Val shape: (1175,) (1175,)
Test shape: (38,) (38,)

```
def sweep_thresholds(probs, labels, num_steps=101):
    thresholds = np.linspace(0.0, 1.0, num_steps)
    rows = []

    for t in thresholds:
        preds = (probs >= t).astype(int) # 1 = blocked, 0 = notblocked

        # metrics for both classes
        precision, recall, f1, _ = precision_recall_fscore_support(
            labels,
            preds,
            labels=[0, 1], # [notblocked, blocked]
```

```

        zero_division=0
    )

    # unpack
    prec_not, rec_not, f1_not = precision[0], recall[0], f1[0]
    prec_blk, rec_blk, f1_blk = precision[1], recall[1], f1[1]

    rows.append({
        "threshold": t,
        "prec_blocked": prec_blk,
        "rec_blocked": rec_blk,
        "f1_blocked": f1_blk,
        "prec_notblocked": prec_not,
        "rec_notblocked": rec_not,
        "f1_notblocked": f1_not,
    })
}

return rows

val_sweep = sweep_thresholds(val_probs, val_labels, num_steps=101)
len(val_sweep), val_sweep[:5]
val_df = pd.DataFrame(val_sweep)
val_df.head()

```

	threshold	prec_blocked	rec_blocked	f1_blocked	prec_notblocked	re
0	0.00	0.566809	1.000000	0.723520	0.000000	
1	0.01	0.566809	1.000000	0.723520	0.000000	
2	0.02	0.573643	1.000000	0.729064	1.000000	
3	0.03	0.592000	1.000000	0.743719	1.000000	
4	0.04	0.630930	0.998498	0.773256	0.991736	

Next steps: [Generate code with val_df](#) [New interactive sheet](#)

Idea 1: pick threshold based on not wanting to miss blockages (high recall), but still trying to keep reasonable precision

```
# filter to thresholds where recall for blocked is >= 0.9
```

```
candidates = val_df[val_df["rec_blocked"] >= 0.9].copy()

if len(candidates) == 0:
    print("No thresholds reach recall >= 0.9 for blocked on val set.")
    # fall back to just the best F1 overall
    best_row = val_df.iloc[val_df["f1_blocked"].idxmax()]
else:
    # among those, pick the one with highest f1_blocked
    best_row = candidates.iloc[candidates["f1_blocked"].idxmax()]

best_threshold = float(best_row["threshold"])
print("Chosen threshold (p_blocked >= t → blocked):", best_threshold)
print(best_row)

# use best_threshold on the test probabilities
def evaluate_with_threshold(probs, labels, threshold):
    preds = (probs >= threshold).astype(int)

    acc = (preds == labels).mean()
    cm = confusion_matrix(labels, preds, labels=[0, 1])
    report = classification_report(
        labels,
        preds,
        labels=[0, 1],
        target_names=["notblocked", "blocked"],
        digits=3
    )
    return acc, cm, report

test_acc_t, test_cm_t, test_report_t = evaluate_with_threshold(
    test_probs, test_labels, best_threshold
)

print(f"Test accuracy at threshold={best_threshold:.3f}: {test_acc_t}")
print("Confusion matrix (rows = true, cols = pred; 0=notblocked,1=blocked)")
print(test_cm_t)
print("\nClassification report:")
print(test_report_t)
```

```
Chosen threshold (p_blocked >= t → blocked): 0.48
threshold          0.480000
prec_blocked      0.917889
rec_blocked       0.939940
f1_blocked        0.928783
prec_notblocked   0.918864
```

```
rec_notblocked      0.889980
f1_notblocked      0.904192
balanced_acc       0.914960
Name: 48, dtype: float64
Test accuracy at threshold=0.480: 0.3684
Confusion matrix (rows = true, cols = pred; 0=notblocked,1=blocked):
[[ 3 24]
 [ 0 11]]

Classification report:
precision    recall   f1-score   support
notblocked    1.000     0.111     0.200      27
    blocked    0.314     1.000     0.478      11

accuracy          0.368      38
macro avg        0.657     0.556     0.339      38
weighted avg     0.802     0.368     0.281      38
```

Idea 2: bc the recall-only approach did worse than the baseline, going to try and change the threshold metric to balanced accuracy ≈ average of class recalls

```
val_df["balanced_acc"] = 0.5 * (val_df["rec_blocked"] + val_df["rec_not_blocked"])

best_row_bal = val_df.iloc[val_df["balanced_acc"].idxmax()]
t_bal = float(best_row_bal["threshold"])
print("Best balanced-accuracy threshold on val:", t_bal)
print(best_row_bal)

test_acc_bal, test_cm_bal, test_report_bal = evaluate_with_threshold(
    test_probs, test_labels, t_bal
)
print(f"\nTest accuracy at threshold={t_bal:.3f}: {test_acc_bal:.4f}")
print("Confusion matrix:")
print(test_cm_bal)
print("\nClassification report:")
print(test_report_bal)
```

Best balanced-accuracy threshold on val: 0.75

threshold 0.750000
prec_blocked 0.956311
rec_blocked 0.887387
f1_blocked 0.920561
prec_notblocked 0.865350
rec_notblocked 0.946955
f1_notblocked 0.904315
balanced_acc 0.917171
Name: 75, dtype: float64

Test accuracy at threshold=0.750: 0.6316

Confusion matrix:

`[[15 12]
 [2 9]]`

Classification report:

	precision	recall	f1-score	support
notblocked	0.882	0.556	0.682	27
blocked	0.429	0.818	0.562	11
accuracy			0.632	38
macro avg	0.655	0.687	0.622	38
weighted avg	0.751	0.632	0.647	38

```
# t=0.5 threshold for comparison
row_05 = val_df.iloc[(val_df["threshold"] - 0.5).abs().idxmin()]
print("Metrics near t=0.5 on val:")
print(row_05)
```

```
Metrics near t=0.5 on val:
threshold          0.500000
prec_blocked       0.920236
rec_blocked        0.935435
f1_blocked         0.927774
prec_notblocked    0.913655
rec_notblocked     0.893910
f1_notblocked      0.903674
balanced_acc        0.914673
Name: 50, dtype: float64
```

Understand why model making mistakes at 0.75 threshold - step 3

```
# using best threshold t_bal from step 2
operating_threshold = float(t_bal)
print("Using best operating threshold (p_blocked):", operating_threshold)

test_ds = test_loader.dataset
print("Test set size:", len(test_ds))

label_id2name = {v: k for k, v in label_map.items()}

# run model on test set and categorize TP / TN / FP / FN
def analyze_test_set(model, dataset, threshold):
    model.eval()
    records = []

    with torch.no_grad():
        for i in range(len(dataset)):
            # dataset.samples should hold (filepath, class_name)
            img_path, label_name = dataset.samples[i]
            pixel_values, label_tensor = dataset[i]

            pixel_values = pixel_values.unsqueeze(0).to(device) # (1,
```

```
        true_id = label_tensor.item()

        logits = model(pixel_values)                      # (1, 2)
        probs = torch.softmax(logits, dim=1)              # (1, 2)
        p_blocked = probs[0, 1].item()                   # prob of cl

        pred_id = int(p_blocked >= threshold)          # 1 = blocked

        # categorize
        if true_id == 1 and pred_id == 1:
            err_type = "TP"    # blocked correctly detected
        elif true_id == 0 and pred_id == 0:
            err_type = "TN"    # notblocked correctly detected
        elif true_id == 0 and pred_id == 1:
            err_type = "FP"    # predicted blocked, actually notblk
        else:   # true_id == 1 and pred_id == 0
            err_type = "FN"    # missed blockage

        records.append({
            "idx": i,
            "path": str(img_path),
            "true_id": true_id,
            "true_name": label_id2name[true_id],
            "pred_id": pred_id,
            "pred_name": label_id2name[pred_id],
            "p_blocked": p_blocked,
            "type": err_type,
        })

    return records

records = analyze_test_set(baseline_model, test_ds, operating_threshold)

TP = [r for r in records if r["type"] == "TP"]
TN = [r for r in records if r["type"] == "TN"]
FP = [r for r in records if r["type"] == "FP"]
FN = [r for r in records if r["type"] == "FN"]

print("\nCounts at threshold", operating_threshold)
print("  TP:", len(TP), "TN:", len(TN), "FP:", len(FP), "FN:", len(FN))
```

Using best operating threshold (p_blocked): 0.75
Test set size: 38

Counts at threshold 0.75

TP: 9 TN: 15 FP: 12 FN: 2

```
# helpers for summarizing and visualizing misclassified
def print_examples(examples, title, max_n=20):
    print(title, f"(showing up to {max_n})")
    for i, r in enumerate(examples[:max_n]):
        print(f"{i+1:2d}. {r['path']}")
        print(f"  true: {r['true_name']:<10s} "
              f"pred: {r['pred_name']:<10s} "
              f"p_blocked={r['p_blocked']:.3f}")
    print()

print_examples(FP, "False Positives (predicted blocked, actually not blocked)")
print_examples(FN, "False Negatives (missed blocked)", max_n=50)

def show_examples(examples, title, n=6):
    n = min(n, len(examples))
    if n == 0:
        print(f"No examples for {title}")
        return

    cols = 3
    rows = (n + cols - 1) // cols

    plt.figure(figsize=(5*cols, 4*rows))
    for i in range(n):
        r = examples[i]
        img = Image.open(r["path"]).convert("RGB")

        plt.subplot(rows, cols, i+1)
        plt.imshow(img)
        plt.axis("off")
        plt.title(
            f"true: {r['true_name']}\n"
            f"pred: {r['pred_name']}\n"
            f"p_blocked={r['p_blocked']:.2f}",
            fontsize=9
        )

    plt.suptitle(title, fontsize=14)
    plt.tight_layout()
    plt.show()

# show a few of each type of mistake
show_examples(FP, "False Positives (not blocked → blocked)", n=8)
```

```
show_examples(FN, "False Negatives (blocked → notblocked)", n=8)
```

False Positives (predicted blocked, actually notblocked) (showing up to 12)

1. /content/split_camera/test/notblocked/2016-09-16 151137 cam68.png
true: notblocked pred: blocked p_blocked=0.811
2. /content/split_camera/test/notblocked/2016-09-16 152433 cam68.png
true: notblocked pred: blocked p_blocked=0.828
3. /content/split_camera/test/notblocked/2016-09-16 151036 cam68.png
true: notblocked pred: blocked p_blocked=0.779
4. /content/split_camera/test/notblocked/2016-09-16 151238 cam68.png
true: notblocked pred: blocked p_blocked=0.827
5. /content/split_camera/test/notblocked/2016-09-16 151726 cam68.png
true: notblocked pred: blocked p_blocked=0.821
6. /content/split_camera/test/notblocked/2016-09-16 152332 cam68.png
true: notblocked pred: blocked p_blocked=0.787
7. /content/split_camera/test/notblocked/2016-09-16 155832 cam68.png
true: notblocked pred: blocked p_blocked=0.932
8. /content/split_camera/test/notblocked/2016-09-16 155731 cam68.png
true: notblocked pred: blocked p_blocked=0.884
9. /content/split_camera/test/notblocked/2016-09-16 155125 cam68.png
true: notblocked pred: blocked p_blocked=0.839
10. /content/split_camera/test/notblocked/2016-09-16 151827 cam68.png
true: notblocked pred: blocked p_blocked=0.771
11. /content/split_camera/test/notblocked/2016-09-16 153105 cam68.png
true: notblocked pred: blocked p_blocked=0.942
12. /content/split_camera/test/notblocked/2016-09-16 152635 cam68.png
true: notblocked pred: blocked p_blocked=0.894

False Negatives (missed blocked) (showing up to 50)

1. /content/split_camera/test/blocked/2016-09-16 153610 cam68.png
true: blocked pred: notblocked p_blocked=0.722
2. /content/split_camera/test/blocked/2016-09-16 153509 cam68.png
true: blocked pred: notblocked p_blocked=0.691





true: notblocked
pred: blocked
p_blocked=0.93



true: notblocked
pred: blocked
p_blocked=0.88



ARLINGTON
VERGINIA



CLARENDON BLVD AT WAYNE



ARLINGTON
VERGINIA

False Negatives

true: blocked
pred: notblocked
p_blocked=0.72



ARLINGTON
VERGINIA



▼ Takeaways

- Model is p good at catching blocked (9/11 --> recall = 0.82)
- We have a high amount of false alarms due to the fact that the lane is very slightly blocked by a truck, but this doesn't count as fully obstructed for a bike to not be able to pass through at all. This isn't bad in my opinion though, because it just further prioritizes biker saftey.
- Model assigns high blocked probabilities (`p_blocked` from 0.78–0.94), showing it may have learned a strong association between visual conetxt cues and the blocked label in training even when the true label is notblocked.
- Only 2 legitimately blocked imgs are missed and both are borderline cases (v close to being predicted as blocked w probabilities 0.722 and 0.691).

Start coding or generate with AI.